

Towards the automatic arrangement of music via quantum annealing

L. Kirby

Durham University

Submitted: 11th January 2025

Music arrangement is usually a complex and time-consuming process; this paper aims to provide an automatic method by which to arrange music via a quantum computing technique called quantum annealing. By splitting a score into a set of phrases, these phrases can form a quadratic unconstrained binary optimisation (QUBO), a function which the quantum computer aims to minimise by choosing the values of discrete variables. At the end of the optimisation process, the resulting chosen values describe the final arrangement, which can then be interpreted into sheet music. This method is tested on an excerpt of Beethoven's String Quartet No. 10 in E-flat major; in this case, the optimisation process is successful, and selects compatible phrases that produce a suitable arrangement.

I. INTRODUCTION

Music is often seen as a very human endeavour. For centuries, only skilled musicians could compose and perform such sequences of sound that would be considered art. Some would go one step further and adapt previously composed pieces for practical or artistic reasons, whether that be in terms of instrumentation, medium, or style, to create an *arrangement*. One of the most famous examples of this is *Pictures at an Exhibition*, a piano suite written by Modest Mussorgsky, but more commonly known for its orchestral arrangement by Maurice Ravel. Traditionally, the arrangement of music is a complex process that requires a deep understanding of musical theory, structure, and style. Composers use their creativity and intuition to create a piece that is both musically interesting and emotionally engaging whilst still remaining faithful to the source material—a challenging and often time-consuming process. Perhaps it is unsurprising, then, that there has been interest in automating this process.

One of the earliest examples of this can be seen in the *Musikalisches Würfelspiel* (“musical dice game”) system popular in the 18th century. The roll of dice would determine the order of pre-composed musical phrases, allowing for the creation of new music without the need for a composer. This system was engaged by the likes of Bach and Mozart, although fell out of fashion the following century. The introduction of computers in the 20th century allowed for more sophisticated methods of music arrangement. Composers could now transpose and manipulate musical parts digitally, without the need to transcribe parts by hand. Moving into the 21st century, more advanced techniques such as genetic algorithms and neural networks have been used to arrange music, with varying degrees of success. However, these methods are limited by the complexity of the problem and the need for extensive training data.

The field of quantum computing has its foundations as early as the 1980s, with the suggestion that hardware following the laws of quantum mechanics could be faster and more powerful than its classical counterpart. Since then, two types of quantum computers have been developed: gate-based and annealing. Gate-based quantum computers, such as those developed by IBM [CITE] and Google [CITE], use quantum gates to manipulate qubits, the quantum equivalent of classical bits, and perform calculations. Quantum annealers, such as those developed by D-Wave Systems, use quantum fluctuations to find the global minimum of a given function. Each are suited to solving different classes of problems, with gate-based computers being

more versatile and annealers being more efficient for certain optimisation problems.

Quantum annealers are particularly well-suited to solving a class of problems known as NP-hard, defined in the field of computational complexity as an optimisation problems where the time taken to find a solution scales exponentially with the input size. A well-known example of an NP-hard problem is the travelling salesman problem, where the goal is to find the shortest route that visits a set of cities exactly once. This problem is difficult to solve because the number of possible routes grows rapidly with the number of cities. Quantum annealers have been shown to be effective at solving NP-hard problems, and have been used to find solutions to a variety of optimisation problems, such as protein folding [CITE], financial portfolio optimisation [CITE], and traffic flow optimisation [CITE].

The combination of quantum computing and music is a relatively new field, but one that has shown promise. Quantum computers have been used to write melodies [1], harmonies [2], create synthesisers [3], and even produce variations of original scores [4], using a mixture of gate-based and annealing methods. However, most of these efforts have been directed at music *composition*, that is, the generation of entirely new sequences and sounds, rather than music *arrangement*, an aspect with very practical use. Quantum computers have the potential to solve complex problems that are intractable for classical computers, and it is proposed that music arrangement is one such problem.

This paper will focus on a subset of music arrangement known as *reduction*—taking a score written for a large number of parts and reducing it to a smaller number of parts, whilst still retaining the musicality and structure of the original. This is a common task in music production, where a piece written for a full orchestra may need to be arranged for a smaller ensemble, such as a string quartet. This forgoes the need to generate new music, as all notes in the arrangement are taken from the original score, and is more readily available to analyse as an optimisation problem.

In this study, the arrangement of a musical score is formulated as an NP-hard graph theory problem, by splitting the score into phrases with associated connections and weightings. This is then solved via a quantum annealer to produce a valid reduction according to imposed constraints. The final arrangement is visualised and evaluated, with suggestions for future work.

II. THEORY

Annealing, in metallurgical terms, is the process of heating and cooling a material to alter its physical properties [5]. The cooling stage is typically slow, allowing particles to arrange themselves into regular lattice structures, making the material softer and less brittle. In the context of quantum computing, quantum annealing is the process of slowly changing a quantum system from one Hamiltonian to another. Much like its metallurgical counterpart, this computational technique allows the system to settle into a more useful final state.

The typical time evolution of a quantum annealer looks like

$$H(t) = \left(1 - \frac{t}{T}\right) H_0 + \frac{t}{T} H_p, \quad (1)$$

where H_0 is the initial ground state that is allowed to evolve to a second state H_p across time T . If this is done adiabatically (in a closed system), with a sufficiently large T , the system will remain in the ground state of $H(t)$, and therefore end in the ground state of H_p at time T . This is the principle of quantum annealing.

The main use of quantum annealing is to solve combinatorial optimisation problems, which are problems that require the minimisation of a function over a discrete set of variables. If H_p is prepared such that its ground state encodes the solution to the optimisation problem, then as long as H_0 and H_p do not commute, the solution will be given at the end of the annealing process.

In order to encode a problem, H_p Hamiltonians take the form of an Ising spin glass, a random arrangement of magnetic dipole moments that can be in one of two states, typically spin-up (+1) or spin-down (-1). A spin glass with a vector s of N spins takes the form

$$H(s) = - \sum_{i < j} J_{ij} s_i s_j - \sum_{i=1}^N h_i s_i \quad (2)$$

where J_{ij} are the coupling strengths, and h_i are the field strengths. The quantum equivalent can be expressed as

$$H_p = H(\sigma^z) \quad (3)$$

where we have replaced the spins with Pauli matrices. This is the *Ising model*; the discrete variables are now *qubits*—binary variables like their classical counterparts, but existing in a superposition of the two states until measurement. The corresponding ground state is prepared with the Hamiltonian

$$H_0 = -h_0 \sum_{i=1}^N \sigma_i^x \quad (4)$$

which is an equal superposition of all possible states in the eigenbasis of H_p [6].

The mapping of a logical problem to an Ising model is called *embedding*. Problems are formulated in an analogous way to a spin glass via a QUBO (quadratic unconstrained binary optimisation). A QUBO takes the form of a function $f(x)$ to be minimised, and is represented by

$$f(x) = \sum_{i < j} Q_{i,j} x_i x_j + \sum_i Q_{i,i} x_i \quad (5)$$

where x is a vector of binary variables (taking the values 0 or 1), and Q is an $N \times N$ upper-diagonal matrix of real weights. The off-diagonal $Q_{i,j}$ terms are known as quadratic coefficients, and diagonal $Q_{i,i}$ terms as linear coefficients. As its name suggests, a QUBO can only include terms up to the second power, restricting coupling to pairwise interactions. This form can easily be transformed to an Ising model using

$$s_i = 2x_i - 1. \quad (6)$$

Once a QUBO has been expressed, it can be sent to a quantum processing unit (QPU) to be solved. The graph architecture of these units allows the mapping of problem QUBOs to physical qubits, with linear terms (x_i) as qubit nodes and quadratic terms ($x_i x_j$) as couplers between them. An embedded system can then be prepared in its initial ground state, and allowed to evolve adiabatically to its final state to obtain the solution.

However, the topology of these graphs often doesn't allow an exact one-to-one mapping; it may be the case that the problem requires a qubit to have more couplers than physically allowed by the QPU. The solution to this is the introduction of chains—a chain of physical qubits can act like a single logical variable, enabling the necessary mapping. The chain strength determines how strongly the chain is coupled, enforcing all qubits within the chain to have the same value. Chain strength is an important parameter that can affect the quality of solutions: if it is too weak, a final state may include chain breaks that return an infeasible solution; if it is too strong, it can overpower the other terms in the QUBO.

The qubits and their couplers define what can be seen as an energy landscape. Different combinations of qubit spins will have different energies, with the lowest energy state being the solution to the problem. A solver (or “sampler”) attempts to characterise the shape of the energy landscape using a random range of input, and returns a number of




Figure 1: A graph of 144 qubits in a D-Wave QPU, using their Pegasus topology. Qubits are represented by nodes, and couplers by edges.

samples with the lowest energy states. Note that the “true” ground state cannot be guaranteed; it could be part of the returned sampleset, or may not even be found at all. It is the parameters of the annealing process that can reduce local minima and increase the likelihood of finding the ground state.

The full embedding process is often handled entirely by the QPU, done “on the fly” each time a problem is submitted. Fixed embeddings can be specified, but these require a priori knowledge of the specific QPU architecture.

III. PROBLEM FORMULATION

Previous approaches to the automatic arrangement of music have relied on classical methods such as neural networks [CITE] or recursive classical algorithms [CITE] in order to produce arrangements that are musically sound. However, these methods are limited by the complexity of the problem and the need for extensive training data. In this study, a new approach to the automatic arrangement of music is proposed by applying quantum annealing.


It has been shown that the arrangement of music can be formulated as a combinatorial optimisation problem: given a set of musical phrases, each phrase can be assigned a binary variable and formulated into a boolean satisfiability (SAT) problem, with variables assigned to clauses according to their compatibility. A valid solution would be an assignment of values such that the boolean expression is satisfied, corresponding to the phrases included in the final arrangement [7]. Phrases are chosen as the smallest unit of music instead of notes to best preserve important melodic lines and harmonies, which may sound dissonant or confusing if split up.

This paper focuses on the reduction of polyphonic (multi-part) music to a single monophonic part, where only one note can be played at a time. A reduction of this form would be relevant for arranging a string quartet for solo violin, for example, whilst still retaining all the important elements from across the instrument parts. This problem can be solved using a reduction of the SAT problem known as the maximum independent set (MIS) problem [6]. Consider a set of musical phrases, each represented by a node x_i in a graph G , with nodes connected by an edge if the phrases overlap. The aim is to find the largest subset of nodes such that no nodes within the subset are connected by an edge—the phrases corresponding to these nodes form the final arrangement. A formulation of this sort ensures that the reduction remains strictly monophonic, as no two chosen phrases play at the same time.


The MIS QUBO takes the form

$$f(x) = A \sum_{ij \in E} x_i x_j - B \sum_i x_i, \quad (7)$$

where we take $x_i = 1$ as the inclusion of that node in the solution set. The first term (penalty term) implements the independence constraint by increasing the energy by A for each edge included, whilst the second term (objective term) is the quantity we are trying to minimise, the (negative) total number of nodes. A and B are known as Lagrange parameters, and can be chosen to balance the two terms accordingly. In this case, the condition $A > B$ must be satisfied, such that any penalty introduced by a node will outweigh its energy reduction.



(a) An example score of three parts, with each phrase coloured.



(b) A graph of the score, with the colour of each node corresponding to its phrase, and edges connecting overlapping phrases.

Figure 2: A toy example demonstrating the construction of the problem graph. It is trivial to see that the maximal independent set contains two nodes.

To increase the quality of the final arrangement, this formulation can further be expanded to include weights on the nodes, which can be used to represent the “musicality” of the phrases. The MIS QUBO then becomes

$$f(x) = A \sum_{ij \in E} x_i x_j - B \sum_i W_i x_i, \quad (8)$$

where W is a vector of real weights. This way, there is bias towards nodes that are considered more musical, and fewer nodes may be chosen as long as it increases the overall perceived quality of the arrangement.

IV. METHODS

D-Wave Systems operates the first and only quantum annealer for commercial access, introduced in 2011 [8]. All problems in this paper are run on the D-Wave Advantage quantum system, which uses their Pegasus architecture, with each qubit connected to 15 others (see Fig. 1). Interaction with the D-Wave QPUs is handled through their Leap quantum cloud service, which is used to submit problems to solvers. D-Wave provides an open-source software development package (SDK) of Python tools which allows users to translate problems into binary optimisation and create QUBOs for the QPU to solve.

Music can be represented by many digital formats, each with their own benefits—this study primarily uses MusicXML [9]. This format focuses on the representation of standard sheet music, describing musical notes, rests, and

other symbols, as well as the structure of the music, such as the arrangement of notes into measures (bars), parts, and scores. It is widely supported by music notation software, allowing translation of music to graphic (PNG, PDF) and audio (MIDI) formats. The Python module `music21` provides extensive resources for manipulating, translating, and creating MusicXML files, allowing scores to be broken down and reconstructed as fit.

The reduction process begins by interpreting a MusicXML file of the original score. Embedded metadata within the score allows it to be broken down into its constituent parts, measures, and notes, as well as containing information about the instrumentation and structure. Each note is represented by a vector of features, such as pitch and duration, as well as its position within the wider piece, which can be measured as an offset (the number of beats from the start of the score).

To construct the problem graph, each individual part needs to be split into a sequence of musical phrases. The approach taken here is the local boundary detection model (LBDM) due to its easy implementation and good performance [10]. The boundaries segmenting a part into phrases are identified using the degree of change between successive notes; larger differences between notes would show an increased likelihood of a boundary, exploiting the fact that the starts and ends of phrases are usually characterised by larger pitch or time intervals [CITE]. Boundary strengths are calculated over two parameters, pitch and inter-onset interval (IOI), which is the difference between the starts of two notes. This is chosen over note duration as it is often more noticeable to the listener. The strength at note i with parameter x_i is given by

$$S_i = x_i \times (r_{i-1,i} + r_{i,i+1}) \quad (9)$$

where $r_{i,i+1}$ is the degree of change of a parameter between notes i and $i + 1$, given by

$$r_{i,i+1} = \frac{|x_i - x_{i+1}|}{x_i + x_{i+1}}. \quad (10)$$

In this way, a note with a high boundary strength would signal the start of a new musical phrase. Each parameter's strength is normalised to the range $[0, 1]$ according to the set of all strengths S via

$$S'_i = \frac{S_i - \min(S)}{\max(S) - \min(S)}. \quad (11)$$

Finally, to find the total boundary strength, the strengths of each parameter are summed with a weighting, using weights derived by trial-and-error (0.3 for pitch and 0.6 for IOI) [10]. If a note's total boundary strength surpasses a specified threshold, it is considered a boundary. Boundaries are always taken at the beginning and end of a score.

Once a list of boundaries is identified, the phrases are extracted by taking all the elements between subsequent boundaries. Each phrase is labelled according to the part it belongs to and its index within that part, allowing the phrases to be easily referenced and reconstructed into a new score once the optimisation is complete. To arrange the phrases into a graph, each phrase is checked against every other phrase for an overlap in time. Phrases are then represented by nodes, with edges connecting phrases that overlap.

The weight of each node is taken here to be the entropy of the phrase. Entropy can be seen as how much information a

variable contains; in this context, a phrase with a higher entropy would indicate a higher level of complexity, increasing the likelihood of it being more "musical". Maximisation of phrase entropy would then result in the creation richer arrangements. For a discrete random variable X with a probability distribution $P(X)$, the Shannon entropy is defined as

$$H(X) := - \sum_i P(x_i) \log_2 P(x_i) \quad (12)$$

where each x_i is a possible value of X and $\sum_i P(x_i) = 1$. In this context, the random variable is a musical note, considering its possible values both in terms of pitch and rhythm.

Pitch entropy is calculated by considering the distribution of pitches in a phrase. The probabilities of each pitch x_i can be found by

$$P(x_i) = \frac{n_i}{N} \quad (13)$$

where n_i is the number of times pitch x_i appears in the phrase and N is the total number of notes in the phrase. A phrase with a greater variety of pitches (and thus greater entropy) will likely be more "interesting" in the final arrangement. Rhythm entropy is calculated in a similar way, but considering the duration of the note instead. This is also calculated using Eq. 13, but instead considering x_i as possible duration values. The total entropy of a phrase is then the sum of the pitch and rhythm entropies.

Once the graph is formed, the weighted MIS QUBO can be constructed via Eq. 8, the Lagrange parameters chosen, and sent to the D-Wave QPU to be solved. The lowest energy solutions returned are sets of nodes that are not connected by edges, whilst maximising their total weighting. The qubits are then matched back to their corresponding phrases, which are reconstructed into a new score, the final reduction.

V. RESULTS

The method was tested on an excerpt of String Quartet No. 10 in E-flat major, Op. 74, by Ludwig van Beethoven, nicknamed the "Harp" quartet. This was chosen due to its relatively simple structure and smaller instrumentation, keeping the problem graph relatively small. Fig. 6a shows the first 13 measures of the original score, taken from a MusicXML file in public domain.

Phrases identified using the LBDM are shown in Fig. 3a, alongside their calculated entropies according to Eq. 12 and their overlaps, represented by edges. A threshold value of 0.4 was used, resulting in a total of 33 identified phrases, with 70 total overlaps. The calculated boundary strengths for the Violin 1 part can be seen in Fig. 9, with the identified boundaries summarised in Tab. I. Many of the nodes form cliques of three or four vertices, that is, each node of the subset is connected to every other node. This can be expected, as often musical phrases across instruments will occur simultaneously, especially in small ensemble music such as this. Notably, there are also two disjoint cliques, subsets that are independent from the main graph. These can be seen as distinct musical ideas that are separated by rests from the other phrases (for example, the first two measures

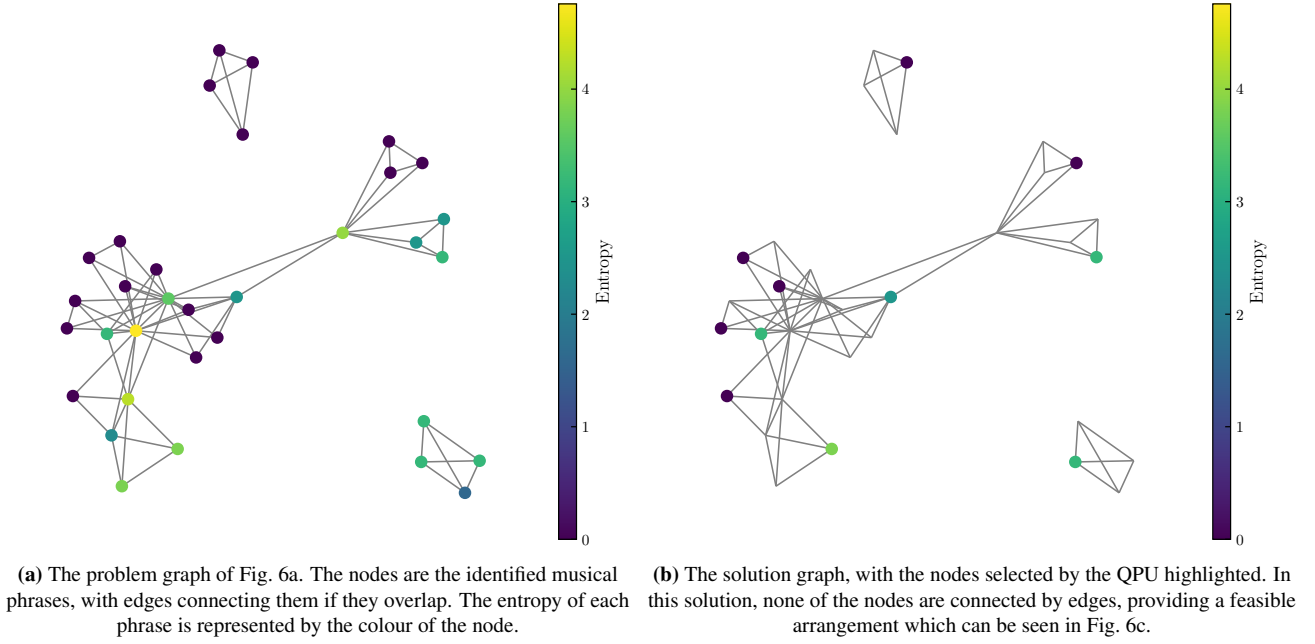


Figure 3: Graph representations of the original score and final arrangement.

of the piece). A non-negligible number of phrases are calculated to have zero entropy—these are phrases identified as only having one note, and thus no variety in pitch or rhythm ($P(x_i) = 1$).

The QUBO was calculated using Eq. 8, with Lagrange parameters $A = 11.5$ and $B = 1$. The value of A was chosen to be twice the maximum node weight (phrase entropy) such that the penalty of choosing a node that introduces an edge into the solution would likely outweigh the benefit of its weighting. After submission of the QUBO to the D-Wave QPU, a sampleset of solutions was returned, which can be seen in Fig. 5. Many solutions had degenerate energies—the lowest energy was -26.8 with a degeneracy of 34 and zero chain breaks. These solutions used slightly different combinations of phrases to make the maximal independent set; this degeneracy is likely due to the high number of cliques, which only allows one node of the subset to be picked, and phrases with zero entropy, which

do not change the final energy when interchanged. One of the lowest energy solutions can be seen in Fig. 3b, with the nodes selected by the QPU highlighted. In this solution, none of the nodes are connected by edges, providing a valid monophonic arrangement. The musical score of this solution is shown in Fig. 6c.




Figure 5: A histogram of the solutions returned by the QPU, over 1000 total reads. Solutions with energies above zero are omitted.




Figure 4: Boundary strengths for each note of the violin part, calculated using Eq. 9. The dashed line shows the threshold value used (0.4), which identifies five total phrases. Note that the beginning and end of the piece are always considered boundaries.

Table I: Summary of the boundary strengths and entropy for the identified phrases of the Violin I part, seen in Fig. 4. The starts of phrases are given by their beat offset from the beginning of the piece.

Start of phrase	Boundary strength	Entropy
0	0	0
2	0.512	3.17
8	0.699	4.00
16	0.611	4.76
36	0.401	3.81

Figure 6 illustrates the process of reducing a string quartet to a single part using quantum annealing. The figure is divided into three main sections: (a) Original excerpt, (b) Selected phrases from solving the MIS problem, and (c) Final arrangement.

(a) Original excerpt: This section shows the beginning of the string quartet, featuring four staves: Violin I, Violin II, Viola, and Violoncello. The tempo is marked "Poco Adagio" and the dynamics are "sotto voce". The music is in A minor, 4/4 time. The excerpt includes measures 6, 10, and 12, with various musical notations such as notes, rests, and dynamic markings like "cresc.", "p", "f", and "espress."

(b) Selected phrases from solving the MIS problem: This section shows the same excerpt as (a), but with selected phrases highlighted in green. These phrases represent the solution to the Maximum Independent Set (MIS) problem, indicating which parts of the original music are retained in the final arrangement. The highlighted phrases are: the first phrase of the Violoncello part (measures 6-8), the first phrase of the Viola part (measures 6-8), the first phrase of the Violin II part (measures 6-8), and the first phrase of the Violin I part (measures 6-8).

(c) Final arrangement: This section shows the final arrangement of the music, which is a single part. It is a reduction of the original quartet, where the selected phrases from (b) are combined into a single melodic line. The final arrangement is in A minor, 4/4 time, and includes measures 6, 10, and 12. The dynamics are "sotto voce" and "espress.".

Figure 6: The beginning of String Quartet No. 15 in A minor, Op. 132, by Ludwig van Beethoven, reduced to a single part.

VI. CONCLUSIONS

In the case of Beethoven’s String Quartet No. 10, this method can be said to be successful in reducing the score to a single monophonic part. A sufficient number of phrases are identified to allow the key melodic lines to be picked out from each part, and the weighted MIS formulation results in a final arrangement that is not musically uninteresting. Here, only one solution was picked arbitrarily; if desired, each of the degenerate lowest energy solutions could be examined individually to isolate the one that is considered most “musical”, safe in the knowledge that all of them are valid.

One advantage this method has over classical algorithms is that no training data is needed. Genetic and deep learning approaches require refinement of their models on large datasets, “teaching” what a valid arrangement looks like, which uses considerable time and resources. By using a QPU, these “rules” can effectively be encoded using constraints, such as the penalty term introduced into the QUBO; as long as these constraints are fulfilled, the QPU will always provide a feasible solution, without any knowledge of previous arrangements.

Another benefit is the solve time: for the problem graph shown in Fig. 3a, the D-Wave QPU access time (the time taken to embed and solve the problem) is approximately 150 ms for 1000 reads. Compared to classical methods using a similar number of phrases, the quantum process is faster by at least one order of magnitude [11]. For classical computers, the time taken to solve NP-hard problems increases exponentially with the problem size, so this speed advantage will only improve as the scores get more complex.

Future work would include testing this process on a wide range of pieces from different musical styles. Scores from different time periods (Baroque vs. Romantic, for example) have wildly different structures, and so the effectiveness of the LBDM could vary in identifying suitable musical phrases. The treatment of melodic lines between different ensembles, whether that be a string quartet or symphony orchestra, also varies, and the isolation of important melodic lines may become more difficult.

To that regard, the effect of LBDM parameters (pitch/IOI weighting, threshold) on phrase identification could be studied. Ideally, musical phrases should be fairly short and similar in length, to give the MIS algorithm the best chance in selecting the most interesting sections of the score. This would prevent important notes being hidden in long phrases with low entropy. Other boundary detection methods could also be tested, with their resulting identified phrases compared qualitatively.

Currently, no regard is taken to the exact instrumentation of the final arrangement. Different instruments have varying ranges of pitches that can physically be played; here, phrases have been chosen from across all the instrument parts (as seen in Fig. 6b), however, this does not mean that the final arrangement can be played by all the instruments as some notes may be unfeasible. Care needs to be taken to check the ambitus (pitch range) of chosen phrases and ensure that it falls within the desired instrument’s range; if not, phrases can be transposed by octaves up or down without altering the melodies significantly.

The MIS formulation works well for reducing a score into a single polyphonic part, as the inclusion of a phrase within the final arrangement lends naturally to the binary variables

used. However, a more useful application would be the ability to reduce a score to any fewer number of parts, whether that be multiple monophonic instruments (e.g. string quartet) or a polyphonic instrument (e.g. piano). To achieve this, the QUBO would need to be altered to allow some edges into the final subset. One suitable formulation would be a graph colouring problem: colouring the vertices of a graph G with a set of n colours such that no edge connects two vertices of the same color. In this context, the number of colours would be the number of desired parts—after solving, each disjoint set of colours would become a monophonic part, the combination of which becomes the final reduction (in this way, the MIS problem can be seen as a colouring problem where $n = 1$).

Although it can be argued that music cannot be objectively “scored”, nonetheless, the quality of the produced arrangements needs to be judged in some way. One suggestion is that the “goodness” of music can only be measured via a Turing-like test, where human subjects are presented with a selection of both human- and computer-generated scores, and asked to categorise them [12]. To this extent, a good measure of the arrangements produced by this method would be to compare them against popular arrangements of the same score composed by human professionals, via a series of blind trials. If the participants fail to identify the computer compositions from the human ones more often than random chance, then it can be said that the generated arrangements are of sufficiently good quality.

To conclude, this paper has shown the automatic reduction of music via quantum annealing. By formulating a score into a function to be minimised, the annealing process can identify parts of the original to become the final arrangement. An automatic process of this sort would be useful to a wide calibre of musicians, removing the time and skill barrier to produce such arrangements, whilst still retaining a sufficient level of quality, keeping music accessible to all. Just as science and business do already, it can be hoped that the arts take advantage of promising new technologies as well.

REFERENCES

- [1] Ashish Arya, Ludmila Botelho, Fabiola Cañete, Dhruvi Kapadia, and Özlem Salehi. Music Composition Using Quantum Annealing, January 2022. arXiv:2201.10557.
- [2] Alex Freedline. Algorithms: Generating Music with D-Wave’s Quantum Annealer, February 2021.
- [3] Eduardo R. Miranda. Quantum Computer: Hello, Music!, June 2020. arXiv:2006.13849.
- [4] Eduardo Reck Miranda and Brian N. Siegelwax. Teaching Qubits to Sing: Mission Impossible?, August 2022. arXiv:2207.08225.
- [5] Oxford English Dictionary. annealing, n., sense 3, December 2024.
- [6] Andrew Lucas. Ising formulations of many NP problems. *Frontiers in Physics*, 2, February 2014. Publisher: Frontiers.
- [7] William S. Moses and Erik D. Demaine. Computational Complexity of Arranging Music, July 2016. arXiv:1607.04220.
- [8] M. W. Johnson, M. H. S. Amin, S. Gildert, T. Lanting, F. Hamze, N. Dickson, R. Harris, A. J. Berkley, J. Johansson, P. Bunyk, E. M. Chapple, C. Enderud, J. P. Hilton, K. Karimi, E. Ladizinsky, N. Ladizinsky, T. Oh, I. Perminov, C. Rich, M. C. Thom, E. Tolkacheva, C. J. S. Truncik, S. Uchaikin, J. Wang, B. Wilson, and G. Rose. Quantum annealing with

- manufactured spins. *Nature*, 473(7346):194–198, May 2011. Publisher: Nature Publishing Group.
- [9] MusicXML for Exchanging Digital Sheet Music.
 - [10] Emiliós Cambouropoulos. The Local Boundary Detection Model (LBDM) and its Application in the Study of Expressive Timing. *International Computer Music Association*, 2011.
 - [11] Jiun-Long Huang, Shih-Chuan Chiu, and Man-Kwan Shan. Towards an automatic music arrangement framework using score reduction. *ACM Trans. Multimedia Comput. Commun. Appl.*, 8(1):8:1–8:23, February 2012.
 - [12] M. Pearce and Geraint A. Wiggins. Towards A Framework for the Evaluation of Machine Compositions. In *Proceedings of the AISB’01 Symposium on Artificial Intelligence and Creativity in the Arts and Sciences*, 2001.
 - [13] Eita Nakamura and Kazuyoshi Yoshii. Statistical piano reduction controlling performance difficulty. *APSIPA Transactions on Signal and Information Processing*, 7:e13, January 2018.
 - [14] Hao-Wen Dong, Chris Donahue, Taylor Berg-Kirkpatrick, and Julian McAuley. Towards Automatic Instrumentation by Learning to Separate Parts in Symbolic Multitrack Music, October 2021. arXiv:2107.05916 [cs].
 - [15] Eduardo Reck Miranda, editor. *Quantum Computer Music: Foundations, Methods and Advanced Concepts*. Springer International Publishing, 2022.
 - [16] Moyu Terao, Eita Nakamura, and Kazuyoshi Yoshii. Neural Band-to-Piano Score Arrangement with Stepless Difficulty Control. In *ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5, June 2023. ISSN: 2379-190X.
 - [17] Yuki Hoshi, Ryohei Orihara, Yuichi Sei, Yasuyuki Tahara, and Akihiko Ohsuga. Versatile Automatic Piano Reduction Generation System by Deep Learning. In *2022 2nd International Conference on Advanced Research in Computing (ICARC)*, pages 66–71, February 2022.

SCIENTIFIC SUMMARY FOR A GENERAL AUDIENCE

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.